

## Challenges



# Overcoming CUDA Compilation Challenges in Cross-Platform Environments

**Authors:** Lino Casu, Akira

## Abstract

The integration of CUDA into diverse environments such as Windows and Linux presents unique challenges, particularly when attempting to create cross-platform applications. This paper explores the difficulties associated with compiling CUDA code on MSYS, WSL2, and native Windows setups, with a focus on achieving compatibility and efficiency across platforms. Practical solutions, including Docker, CMake, and hybrid toolchains, are proposed to address these issues.

---

## 1. Introduction

CUDA has become a cornerstone for high-performance parallel computing, enabling significant speedups in tasks such as numerical simulations and large-scale mathematical computations. Despite its potential, compiling CUDA code in heterogeneous environments introduces a host of challenges:

1. **Toolchain Mismatch:** CUDA's tight integration with specific compilers (e.g., Visual Studio on Windows) limits flexibility.
  2. **Cross-Platform Compatibility:** Creating portable binaries for both Linux and Windows is non-trivial.
  3. **Environment Constraints:** Tools like MSYS and WSL2 impose unique restrictions on file systems and permissions.
- 

## 2. Challenges in CUDA Compilation

### 2.1 MSYS on Windows

MSYS, while providing a Unix-like environment on Windows, is not designed to handle native CUDA toolchains:

- **Path Translation Issues:** CUDA's `nvcc` expects native Windows paths, conflicting with MSYS's Unix-style paths.
- **Permission Errors:** File system emulation in MSYS introduces inconsistencies in access control.

- **Lack of Direct Support:** NVIDIA does not officially support MSYS for CUDA development.

## 2.2 WSL2 with CUDA

Windows Subsystem for Linux (WSL2) adds a Linux kernel layer to Windows, supporting CUDA since 2020. However, challenges remain:

- **Driver Requirements:** WSL2 requires specific versions of NVIDIA drivers.
- **Cross-Compilation Complexity:** Producing Windows executables from a Linux environment remains cumbersome.
- **Performance Overheads:** Although WSL2 is efficient, there is a slight latency due to virtualization.

## 2.3 Native Windows Toolchains

Using native Windows tools, such as Visual Studio with CUDA, provides robust support but introduces limitations:

- **Closed Ecosystem:** Integration with non-Microsoft tools like MinGW or MSYS is difficult.
- **Portability Issues:** Executables generated on Windows may not work seamlessly in Linux environments.
- **Build Automation:** Automating builds across multiple platforms requires additional setup.

---

## 3. Proposed Solutions

### 3.1 Using WSL2 for Development

- **Advantages:**
  - Native Linux environment for CUDA development.
  - Direct access to GPU resources with minimal configuration.
- **Implementation Steps:**
  1. Install WSL2 and ensure NVIDIA drivers are up-to-date.
  2. Install the CUDA Toolkit for Linux.
  3. Develop and test within WSL2, then cross-compile using tools like MinGW if needed.

### 3.2 Native Windows with CMake

- **Advantages:**
  - Unified build system for cross-platform development.
  - Direct integration with CUDA and Visual Studio.
- **Implementation Steps:**
  1. Configure a CMakeLists.txt file with CUDA support.
  2. Use CMake's generator feature to create project files for both Windows and Linux.
  3. Compile using nvcc and target static linking for portability.

### 3.3 Portable Docker Environments

- **Advantages:**
  - Encapsulation of all dependencies and toolchains.
  - Eliminates environment-specific issues.
- **Implementation Steps:**
  1. Create a Dockerfile with the necessary CUDA, GMP, and other dependencies.
  2. Mount source code into the container.
  3. Compile and test within the container, ensuring compatibility with the host system.

### 3.4 Statically Linked Executables

- **Advantages:**
  - Single binary that runs across compatible systems.
  - Reduces runtime dependency issues.
- **Implementation Steps:**
  1. Compile with static linking using flags like -static or -lgmp.
  2. Test for compatibility on both Windows and Linux.

---

## 4. Implementation Example: CMake with CUDA

### Sample CMakeLists.txt for Cross-Platform Builds

```
cmake_minimum_required(VERSION 3.18)
```

```
project(CUDA_CrossPlatform LANGUAGES CXX CUDA)
```

```
set(CMAKE_CUDA_STANDARD 14)
```

```
add_executable(cuda_app main.cu)
```

```
set_target_properties(cuda_app PROPERTIES
```

```
    CUDA_SEPARABLE_COMPILATION ON
```

```
    CUDA_ARCHITECTURES 75)
```

```
if(WIN32)
```

```
    target_link_libraries(cuda_app PRIVATE gmp gmpxx)
```

```
else()
```

```
    target_link_libraries(cuda_app PRIVATE -static gmp gmpxx)
```

```
endif()
```

- **Steps to Build:**

- On Windows: Run `cmake . -G "Visual Studio 16 2019" && cmake --build ..`
- On Linux: Run `cmake . && make`.

---

## 5. Discussion

### Trade-offs Between Approaches

Approach	Pros	Cons
WSL2	Linux-native tools, GPU support	Requires updated drivers
Native Windows (VS)	Full CUDA support	Limited portability
Docker	Fully portable environments	Higher setup complexity
Static Linking	Portable binaries	Larger executable sizes

### Recommendations

For development, WSL2 offers the best balance of performance and compatibility. For deployment, Docker ensures consistent environments across systems.

---

## **6. Conclusion**

The challenges of CUDA compilation in cross-platform environments can be mitigated through a combination of modern tools and best practices. By leveraging WSL2, Docker, and CMake, developers can create robust and portable solutions for high-performance computing tasks.

## **Future Work**

- Integration of Multi-GPU scaling in Docker.
- Automation scripts for end-to-end builds.
- Further optimization of CUDA kernels for hybrid environments.

---

## **References**

1. NVIDIA CUDA Toolkit Documentation.
2. Docker: Building CUDA Images.
3. GNU MP: The GNU Multiple Precision Arithmetic Library.
4. CMake: Cross-Platform Makefile Generator.